

Statistics for Data Analytics


Summary Sheet

John S Butler (TU Dublin)


Measures of Location

Different aspects of a distribution of data can be summarised by the three measures of location:

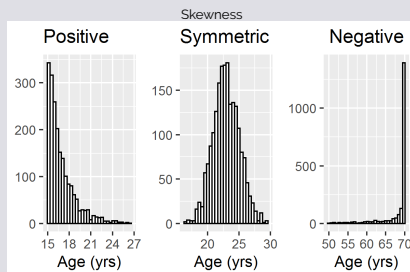
First Moment: Middle



Second Moment: Spread



Third Moment: Symmetry




Left: One Direction Concert Attendance, Middle: Harry Styles Concert Attendance, Right: André Rieu Concert Attendance

Data Type

- Categorical
- Ordinal
- Interval
- Ratio

But a picture is worth having



Mathematical Probability

Definitions

Define some event A that can be the outcome of an experiment. $\Pr(A)$ is the probability of a given event A that will happen.

Rules:

- $\Pr(A)$ is between 0 and 1. $0 \leq \Pr(A) \leq 1$;
- $\Pr(A) = 1$, means it will definitely happen;
- $\Pr(A) = 0$, means it will definitely **not** happen;
- $\Pr(A) = 0.05$, is arbitrarily considered unlikely.

Sample Space and Events

The **Sample Space**, S , of an experiment is the universal set of all possible outcomes for that experiment, defined so, no two outcomes can occur simultaneously. For example:

- Throwing a die $S = \{1, 2, 3, 4, 5, 6\}$.

An event, A , is a subset of the sample space S . For example:

- Throwing a die $S = \{3, 4, 6\}$;

Axioms of Probabilities

For an event A subset S associated a number $\Pr(A)$, the probability of A , which must have the following properties

- $\Pr(A \cap B) = 0$; $\Pr(A \cup B) = \Pr(A) + \Pr(B)$;
- Probability of the Null Event $\Pr(\emptyset) = 0$;
- The probability of the complement of A , $\Pr(\bar{A}) = 1 - \Pr(A)$;
- $\Pr(A \cup B) = \Pr(A) + \Pr(B) - \Pr(A \cap B)$.

Conditional Probability

The Conditional Probability $\Pr(A|B)$ denotes the probability of the event A occurring given that the event B has occurred,

$$\Pr(A|B) = \frac{\Pr(A \cap B)}{\Pr(B)}.$$

Example: The rain in Ireland

A normal probability would be what is the probability it is going to rain, $\Pr(\text{rain})$. A conditional probability would, be what is the probability it is going to rain **given** that you are in Ireland, $\Pr(\text{rain}|\text{Ireland})$,

$$\Pr(\text{rain}|\text{Ireland}) = \frac{\Pr(\text{rain} \cap \text{Ireland})}{\Pr(\text{Ireland})},$$

where the probability of rain is $\Pr(\text{rain}) = 0.3$, the probability of being in Ireland is $\Pr(\text{Ireland}) = 0.4$ and the probability of being in Ireland and it raining is $\Pr(\text{rain} \cap \text{Ireland}) = 0.2$.

$$\Pr(\text{rain}|\text{Ireland}) = \frac{0.2}{0.4} = 0.5,$$

You could be interested in the probability that you are in Ireland **given** that it is raining,

$$\Pr(\text{Ireland}|\text{rain}) = \frac{\Pr(\text{rain} \cap \text{Ireland})}{\Pr(\text{rain})} = \frac{0.2}{0.3} = 0.75.$$

Bayes Theorem

Bayes Theorem states

$$\Pr(A|B) = \frac{\Pr(B|A)\Pr(A)}{\Pr(B)}.$$

Example: Diagnostic test

The probability that an individual has a rare disease is $\Pr(\text{Disease}) = 0.01$. The probability that a diagnostic test results in a positive (+) test *given you have* the disease is $\Pr(+|\text{Disease}) = 0.95$. On the other hand, the probability that the diagnostic test results in a positive (+) test *given you do not have* the disease is $\Pr(+|\text{No Disease}) = 0.1$. This raises the important question if you are given a positive diagnosis, what is the probability you have the disease $\Pr(\text{Disease}|+)$? From Bayes Theorem we have:

$$\Pr(\text{Disease}|+) = \frac{\Pr(+|\text{Disease}) \Pr(\text{Disease})}{\Pr(+)}$$

The probability of a positive test is,

$$\Pr(+) = \Pr(+|\text{Disease}) \Pr(\text{Disease}) + \Pr(+|\text{No Disease}) \Pr(\text{No Disease}),$$

$$\Pr(+) = 0.1085.$$

$$\Pr(\text{Disease}|+) = \frac{\Pr(+|\text{Disease}) \Pr(\text{Disease})}{\Pr(+)} = \frac{0.95 \times 0.01}{0.1085} = 0.0875576.$$

This can also be done in a simple table format, by assume a population of 10,000

Group	+ Diagnosis	- Diagnosis	Total
Disease	95	5	100
No Disease	990	8,910	9,900
Total	1,085	8,915	10,000

From the table we can calculate the same answer,

$$\Pr(\text{Disease}|+) = \frac{95}{1085} = 0.0875576.$$

Discrete Distribution

Probability Mass Functions

Event Number i	0	1	2	3	4
Event Value x_i	-1	0	1	2	3
Probability of Event $\Pr(x_i)$	0.3	0.1	0.3	0.1	0.2

The expected value of the distribution is:

$$\mu = E[X] = \sum_i x_i \Pr(x_i),$$

$$\sum_i x_i \Pr(x_i) = -1 \times 0.3 + 0 \times 0.1 + 1 \times 0.3 + 0.1 \times 2 + 0.2 \times 3 = 0.8,$$

The variance of the distribution is:

$$\text{Var}[X] = \sum_i (x_i - \mu)^2 \Pr(x_i) = \sum_i (x_i - 0.8)^2 \Pr(x_i).$$


Binomial Distribution

The formula for the Binomial distribution is:

$$\Pr(k) = \binom{n}{k} p^k q^{n-k}, \quad k = 0, 1, 2, \dots, n,$$

$$E[k] = np, \quad \text{Var}[k] = npq,$$

where n is the total of games, k is the number of 'wins', p is the probability of a 'win', $q = 1 - p$ probability of a 'loss'.




Geometric Distribution

The formula for the Geometric distribution is:

$$\Pr(k) = q^{(k-1)} p, \quad k = 1, 2, \dots$$

$$E[k] = \frac{1}{p}, \quad \text{Var}[k] = \frac{q}{p^2},$$

k is the number of events until one 'win', p is the probability of a 'win', $q = 1 - p$ probability of a 'loss'.



Discrete Distribution


Poisson Distribution

The formula for the Poisson distribution is:

$$\Pr(k) = \frac{\lambda^k e^{-\lambda}}{k!}, \quad k = 0, 1, 2, \dots$$


$$E[k] = \lambda, \quad \text{Var}[k] = \lambda.$$

where λ is the mean and standard deviation of the distribution and k is the number of 'wins' in a specified time or space.




Continuous Distribution

Normal Distribution



Confidence Intervals



Hypothesis Testing

Five steps for Hypothesis testings

1. State the Null Hypothesis H_0 ;
2. State an Alternative Hypothesis H_a ;
3. Calculate a Test Statistic (see below);
4. Calculate a p-value and/or set a rejection region;
5. State your conclusions.

z-test

Continuous Data

The test statistic is given by

$$Z = \frac{\bar{x} - \mu}{\frac{\sigma}{\sqrt{n}}} \sim \mathcal{N}(0, 1),$$

where \bar{x} is the observed mean, μ is the historical mean, σ is the standard deviation and n is the number of observations. $\mathcal{N}(0, 1)$ is the normal distribution with a mean of 0 and a standard deviation of 1.

Do supplements make you faster?

The effect of a food supplements on the response time in rats is of interest to a biologist. They have established that the normal response time of rats is $\mu_0 = 1.2$ seconds. The $n = 100$ rats were given a new food supplements. The following summary statistics were recorded from the data $\bar{x} = 1.05$ and $\sigma = 0.5$ seconds

1. The rats in the study are the same as normal rats, $H_0 : \mu = 1.2$.
2. The rats are different, $H_a : \mu \neq 1.2$.
3. Calculate a Test Statistic $Z = \frac{1.05 - 1.2}{\frac{0.5}{\sqrt{100}}} = -3$
4. Reject the Null hypothesis H_0 if $Z < -1.96$ and $Z > 1.96$
5. The data suggests that rats are faster with the new food.

Proportional Data

The test statistic is given by

$$z = \frac{\hat{p} - p_0}{\sqrt{\frac{p_0 q_0}{n}}} \sim \mathcal{N}(0, 1).$$

where \hat{p} is the observed proportion, p_0 is the historical proportion, q_0 is the complement $q_0 = 1 - p_0$, and n is the number of observations.

t-test

paired t-test

The test statistic is given by

$$t = \frac{\bar{x} - \mu_0}{\frac{s}{\sqrt{n}}} \sim t_{\alpha, df}$$

where \bar{x} is the observed mean, μ_0 is the null mean, s is the standard deviation and n is the number of observations. α is the alpha level and df is the degrees of freedom.

unpaired t-test

The test statistic is given by

$$t = \frac{\bar{x}_1 - \bar{x}_2}{s_p \sqrt{\frac{1}{n_1} + \frac{1}{n_2}}} \sim t_{\alpha, df}$$

where $s_p = \sqrt{\frac{s_{x_1}^2 + s_{x_2}^2}{2}}$ is the pooled sample standard deviation, \bar{x}_1 and \bar{x}_2 are the sample means, n_1 and n_2 are the sample sizes.

χ^2 Independence test

The test statistic to test if data are independent of group is given by:

$$\chi_{Ind}^2 = \sum \frac{(O - E)^2}{E} \sim \chi_{(r-1)(c-1)}^2$$

where O is the observed data, E is the expected data if independent, r is the number of rows and c is the number of columns.

Does ice-cream flavour matter?

An ice-cream company had 500 people sample one of three different ice-cream flavours and asked them to say whether they liked or disliked the ice-cream.

	Vanilla	Chocolate	Strawberry
Liked	130	170	100
Disliked	20	30	50

The χ_{Ind}^2 independence test could be used to determine if the enjoyment of the ice-cream depends on the flavour.

χ^2 Goodness of Fit


The test statistic to test if data come from a specific distribution is given by:

$$\chi_{GoF}^2 = \sum \frac{(O - E)^2}{E} \sim \chi_{k-1}^2$$

where O is the observed data, E is the expected data from a chosen distribution and k is the number of observation bins.

Does it fit?

The χ_{GoF}^2 can test if the observed distribution of the height of Dutch people (grey) fits the expected distribution of heights (dark grey).



Linear Regression

A linear regression is used to model a linear relationship of the dependent variable y and the regressors x_1, x_2, \dots

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots$$


where β_0, β_1 are the slopes of the regressors.

Height Prediction

A simple linear regression (correlation) is used to predict the height of 744 children y using the height of their parent x .

$$y = \beta_0 + \beta_1 x$$

The plot below shows the fit of the model:



The parents' height x explained 12.7% of the childrens' height y .

Logistic Regression

A logistic regression (or logit model) is used to model the probability of a binary events such as win/lose. The general formula for the Logistic regression is

$$p_i = \frac{e^\eta}{1 + e^\eta}$$

where

$$\eta = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots$$

and β is the slope corresponding to the predictor variable x .

Sexton Conversion Rate


Data from 1000 conversions kicks by Johnny Sexton was acquired; the distance (m) from the goal-line and if the kick was a miss 0 or a conversion 1. The data was fit to a logistic regression. The model was

$$p = \frac{e^\eta}{1 + e^\eta}$$

where

$$\eta = \beta_0 + \beta_1 \text{Distance}$$

and p is the probability of a conversion. The plot below shows the fit of the model:



The model predicts that at the half-way line (50m) Sexton has a 0.375 probability of conversion.

Bibliography

- Alexander, R. - Telling Stories with Data 2022 [website](#)
- Devore & Peck - Statistics: The exploration and analysis of data (2011)
- James, G., Witten, D., Hastie, T., & Tibshirani, R. (2013). An introduction to statistical learning (Vol. 112, p. 18). New York: springer [book](#) [website](#).
- Poldrack R. Statistical Thinking in the 21st Century 2020 [website](#).

Popular Press

Fry, H. - Hello World: How to be Human in the Age of the Machine, Doubleday, 2018

Resources

Butler, J. S., [R GitHub Repository](#)

Notation

- \bar{x} - mean of a list of numbers x_i
- σ - standard deviation of a list of numbers x_i
- σ^2 - variance of a list of numbers
- $\Pr(A)$ - probability of event A
- $\Pr(\bar{A})$ - probability of not event A
- $\Pr(A|B)$ - probability of event A given event B is known
- $\sum_i^n x_i$ - the sum of a list of number x_i
- $n!$ - n factorial is $n \times (n-1) \times \dots \times 1$
- $5!$ - 5 factorial is $5 \times (5-1) \times (5-2) \times (5-3) \times (5-4) = 5 \times 4 \times 3 \times 2 \times 1 = 120$
- $\binom{n}{k} = {}^n C_k$ - n choose k equals to $\frac{n!}{k!(n-k)!}$
- $\binom{5}{3} = {}^5 C_3$ - 5 choose 3 equals to $\frac{5!}{3!(5-3)!} = \frac{5!}{3!2!} = \frac{5 \times 4 \times 3 \times 2 \times 1}{3 \times 2 \times 1 \times 2 \times 1} = 10$
- ${}^n P_k$ - n pick k equals to $\frac{n!}{(n-k)!}$
- ${}^5 P_3$ - 5 pick 3 equals to $\frac{5!}{(5-3)!} = \frac{5!}{2!} = \frac{5 \times 4 \times 3 \times 2 \times 1}{2 \times 1} = 60$
- p - p probability of a "win"
- q - q probability of a "loss" $1 - p$
- p^n - p to the power of n is $p \times p \times \dots \times p$
- 0.1^4 - 0.1 to the power of 4 is $0.1 \times 0.1 \times 0.1 \times 0.1$
- $E[X]$ - the expected value of a probability distribution
- $\text{Var}[X]$ - the variance of a probability distribution
- e - is the exponential which is it equal to approximately 2.718 it is comes up again and again in mathematics formulas
- H_0 - null hypothesis
- H_α - alternative hypothesis
- μ - real mean (generally never known)
- μ_0 - historical mean
- \bar{x} - observed mean given the data
- \hat{p} - is the observed sample proportion
- p_0 - is the historical proportion
- $\mathcal{N}(\mu, \sigma)$ - is the Gaussian distribution with mean μ and standard deviation σ
- $\mathcal{N}(0, 1)$ - is a special case of Gaussian distribution known as the Normal Distribution with mean 0 and standard deviation 1
- df-degrees of freedom
- χ_{df}^2 - Chi (χ)-squared (2) distribution with degrees of freedom df
- β the coefficient for a regression
- $\hat{\beta}$ the coefficient estimated for a regression from the observed